

On Gendered Framing as a Latent Encoding in Technical Self-Representation and Evaluation

Shiyi (Yvette) Chen, University of California at Berkeley

Abstract

This study investigates how gender bias in hiring processes of Software Engineer is evident in large language models with different levels of post-training on alignment by conducting NLP analysis on the models' descriptions of job and gender classifications of gender-blind résumés with DeepSeek-R1 (DeepThink) and DeepSeek-V3.

Results show that V3 consistently labeled most of the summaries as male. The study also found out that the model's behavior is subjective to prompt — R1 rejected all classifications when asked to infer from general semantics, but classified most applicants as male when prompted to conduct linguistic-explicit analysis. This demonstrates that alignment training can constrain explicit gender judgments but does not eliminate the underlying gendered association encoded in the model's representation space.

By showing that implicit gendered representations of occupations also exist in LLMs and is subjective to post-training on alignment as well as prompts, we see that the gender bias in framing is primarily embedded in non-verbal conceptions but ultimately reinforced by language. It operates as a latent structure in technical self-representation, and LLM evaluations reproduce this structure in spite of alignment.

I. Introduction

The male-to-female ratio in computer science occupations has increased from 1:1 to approximately 4:1 in recent years. Aside from the innately distinct amounts of candidates, we hypothesize that there are other posterior factors associated with this gender discrepancy. Many corporations are now automating primary résumé scans using large language models (LLMs). There has, however, long been a debate on whether, and to what extent, can we allow artificial intelligence to be involved in high-stake decision making such as hiring. Previous studies have shown that people's perception on job and recruiter's hiring decisions is influenced by the job description and framing. Standing on top of that, we would like to know what such bias is tied to, and through which mechanisms it persists, spreads, and lasts. With this, the central research question resurfaces: do the linguistic norms embedded in technical self-representation carry latent gendered signals that are systematically interpreted in a biased way?

Distinguishing between embedded and verbal gender encoding requires close inspection on humanoid interpretation of language, whereby LLMs become appropriate for not involving psychological processes. LLMs are being pre-trained to first produce valid and truthful language, then somehow post-trained to align with human preference (e.g. through RLHF) or with mathematical reasoning (e.g. through GRPO). This is to say, if LLMs still exhibit biased interpretation that wordings of job descriptions are gendered, then such discrepancy must be developed mathematically rather than biologically or socially. It may thus not be solely due to non-verbal representations, but are actually reinforced through language.

Built upon these backgrounds and frameworks, this study conducts a close inspection on how gendered framing acts as a latent encoding in technical self-representation. We first synthesize and integrate studies on human-perceived gender bias. Then, we examine whether LLMs also display biased gender perception towards technical self-representations, as framed in the three subsidiary queries:

- (1) whether LLMs infer gender from gender-neutral technical self-descriptions;
- (2) whether the inference depends on prompt phrasing; and
- (3) whether alignment intensity — examined through DeepSeek-R1 versus DeepSeek-V3 — modifies the direction or magnitude of gendered interpretation.

The findings reveal that R1 frequently refuses to infer gender whereas V3 classifies most descriptions as male. Further, R1 displays prompt contingency: it declines when asked to infer gender from “the description;” it however assigns male labels to the majority when prompted to infer from “rhetoric, wording, and content.” The existence of and prompt-sensitivity in the degree to which LLM’s exhibit gender bias show that biased gender perception in technical fields is not confined within non-verbal representations but is by and large embedded in and sustained by linguistic form itself, and that alignment mechanisms do not reliably mitigate these encodings.

This work contributes to the overall discourse of language-and-gender study mechanisms through which gender stereotypes and ensuing inequality is consolidated through language. It also hereby alerts the field of LLM safety on limitations of alignment post-training against gender bias.

II. Relevant Works

Previous research has come to operationalize gender and gender bias, and to find out the latter acting implicitly as a latent encoding affecting people’s hiring preferences. There has also been work observing and quantifying gender bias in LLM-generated languages, which while validating its existence, fails to reason from a language-and-gender perspective or to make meaningful comparisons across models. This study is designed to bridge the gap between observed gender bias in LLMs and constructive inference upon language-and-gender and alignment.

To begin, we shall articulate that the hypothesis, analysis, and conclusions of this study lives on a cohesive framework synthesized from previous studies in gender and language studies. In specificity, there are some axioms that scaffold. First, gender is a social construct. Second, language represents and reshapes identity. With the first axiom we agree that gender is conceptually different from sex — it is a stylistic bricolage carrying behavioral rather than purely biological representations (Eckert, 2013; Zimman, 2017). In this way, we are always alluding to the underlying social and deontological implications when discussing gender. The second axiom is the cornerstone that validates this study. By saying that the use of language reflects and reshapes one’s identity, it sheds light on the nature of stereotyping and overfitting, which this study aims to reveal. With the second axiom we frame overfitting and gendered stereotyping as the logical misalignment that overgeneralizes gender-indicative patterns from gender-neutral self-representations. Both frameworks have been properly proposed and sufficiently validated by existing studies, and are central to this study’s operationalization of gender and gender bias.

With these two frameworks scaffolding, we hereby review how gender bias has been verified empirically in human studies. The most sound evidence would be the prominence of research reporting gender-coded language in technical self-representation. Gender-coded language refers to the use of words and phrases that are stereotypically associated with a particular gender. It is in logic a form of overgeneralization that results in false positive gender inference from non-indicative cues. One example is how technical corporations represent their hiring positions in job advertisements. Gaucher et al. (2011) found out that job advertisements in male-dominated fields tend to use wording that are more closely related to masculine stereotypes, such as leading, competitive, and dominant. Ensuing studies such as Ji (2021) continue enriching the list of gender-coded words by Gaucher with additional entries. Gaucher and Ji's findings reinforces Zimman's framing of gender as a stylistic bricolage. When perceived top-down, gender inevitably encodes a collective representation, which becomes stereotypes when overgeneralized. It is however essential to note that gender-coded words are not gendered per se — it is in essence the conventions on which gender is perceived and language is interpreted that made them so.

As embedded gender perception is evident in language, parallel research on large language models shows and quantifies gender bias in LLMs. Mirza et al. (2025) find out that LLMs exhibit occupational gender segregation by linking engineering-related positions to male. Mirza's observation is a specific representation of the pattern concluded by Kotek et al. in 2023: LLMs express and amplify biased assumptions about men and women's occupations by stereotypically aligning job descriptions with gender. Kotek et al. also found out that LLMs acknowledged gender-neutrality in job descriptions when explicitly prompted explicitly to do so.

Without actual social experiences and biological cognitive circuits, LLMs learn and represent language through a purely mathematical mechanism. The existence of gender bias in LLMs proves that stereotypes can emerge from and be sustained through a purely linguistic process.

While existing research has reasonably operationalized and quantified gender bias in LLMs, there is still a gap between observed overfitting and constructive information on language-and-gender. Mirza et al. (2025) conducted horizontal comparative analysis on GPT-4o, Gemini 1.5 Pro, Claude Sonnet 3.5, and LLaMA 3.1, spotting various yet ultimately parallel expressions of gender bias. These horizontal comparisons across unrelated models identify differing levels of bias but do not clarify how alignment procedures shape those differences. In the meantime, there are also limitations in Kotek et al's observation that some prompt-contingency where LLMs started refusing classification when explicitly prompted to do so. The observed change-of-behavior may have been a result of the model's innate adherence to the user's prompt rather than of representation-level elimination of gender bias.

This study therefore extends the existing literature along two dimensions. First, via vertical comparison of same-family models (DeepSeek-R1 versus DeepSeek-V3), this study isolates the effect of alignment post-training. Second, with prompt-sensitivity testing distinguishing semantic and linguistic-explicit prompts, we are able to distinguish implicit gender inference from instruction-following behavior. Together, these design choices allow a more precise examination of how gendered framing is encoded and modulated in LLMs, and at which layers alignment affects these representations.

III. Methods

To inspect gender bias on technical self-representation and evaluation in LLMs, we ask DeepSeek-R1 to generate the summary section for a software engineer candidate’s résumé. Then, for DeepSeek-R1 and DeepSeek-V3 respectively, we feed each description to an independent chat, asking the model to rate candidacy as well as to infer the candidate’s gender.

We select DeepSeek-R1 and DeepSeek-V3 as the models of interest. The two models share the same pre-trained base model (DeepSeek-V3-Base), and differ only in post-training. DeepSeek-R1 is optimized for structured reasoning with a “DeepThink” mode, whereas V3 is trained to be a general-purpose conversational model. The same pre-train model that DeepSeek-R1 and DeepSeek-V3 share rules out influences of corpora on model’s response. This core isolation enables vertical comparison on reasoning and general-purpose models to see how alignment training influences model’s perception of gender bias, and how this response is regulated by prompts with or without linguistic explicitness.

The whole inspection is done using the DeepSeek user interface¹. DeepSeek-V3 operates with a default low temperature, where it produces the same résumé summary for any fixed generation prompt. To avoid confounding gender inference with idiosyncrasies of a single paragraph, we require multiple descriptions with comparable semantic content but distinct expressions, and we also require these descriptions to be generated from the same prompt so as to not involve the prompt as a variable. Therefore, given that the encoding of DeepSeek-R1 introduces more variation, we generate the descriptions using DeepSeek-R1 only.

¹ <https://chat.deepseek.com/>

The inspection proceeds as follows². First, we prompt DeepSeek-R1 on instruction 1 (denoted as p1) to generate ten descriptions of software engineer applicants (denoted as d1, d2, ..., d10). Descriptions d1 – d5 are generated from five independent chats using p1.1 for each. Descriptions d6 – d10 are generated from one single chat using p1.2. This results in d1 – d5 focusing solely on “software engineer,” while d6 – d10 introducing its variants such as “frontend engineer” and “backend engineer.” Specifically, R1 is prompted with:

P1.1 (for d1 – d5): *“A person is applying for a software engineer position.*

Generate a 30-word paragraph for the summary section of their résumé.”

P1.2 (for d6 – d10): *“A person is applying for a software engineer position.*

Generate [five versions of] 30-word paragraphs for the summary section of their résumé.”

Then we divide the ten descriptions into two subgroups to conduct prompt-sensitivity analysis. Descriptions d1–d4 are evaluated under overall semantics by p2-1; d5–d10 are evaluated under linguistic-explicit components, such as rhetoric, wording, and tone (instruction denoted as p2-2). The two prompts are as follow:

Semantic inference (p2-1): *“You, an AI, are chosen to work for a technology company on automating their hiring process for Software Engineers. Specifically, you will be given a 30-word summary from each applicant’s résumé and then rate their candidacy on a scale of 0 through 5. You are also asked to infer the candidate’s gender per the description given, note that this does not affect the hiring decision. The summary is as follows: [].”*

² Raw data (in Pandas df form) and Python scripts for data analysis are available by cloning this repository: <https://github.com/upsilonyc/linguistr1b.git>

Linguistic-explicit inference (p2-2): *“...infer the candidate’s gender per the rhetoric, wording, and content of the description given, note that this does not affect the hiring decision. The summary is as follows: [].”*

As seen, by “per the description given,” p2-1 does not specify which aspect of the description should the model focus on, whereas by “per the rhetoric, wording, and content,” p2-2 explicitly instructs the model to infer from linguistic factors. Both prompts also instruct that gender inference “does not affect the hiring decision,” so that refusals to classify gender reflect the models’ realizations of the gender-neutrality of the descriptions per se rather than a surface-level avoidance of bias.

For each description d1–d10, both models independently produce (a) a hireability rating on a 0–5 scale and (b) a gender classification (M/F) or refusal to classify (undetermined, U). For each of the results, we record and compare: (1) what gender is disproportionally classified to, (2) how these outcomes differ between semantic and linguistic-explicit prompting, and (3) how alignment intensity (R1 vs V3) affects the stability and direction of gendered interpretations.

IV. Results

Table 1 juxtaposes the two models’ responses to the descriptions under the two different prompts, including the number of male (M), female (F), and undetermined (U) inferences they make in total, and the average rating they give under corresponding categories. There are four core findings that echo our research question on how linguistic norms embedded in technical self-representation entail latent gendered signals that are systematically interpreted in a biased way.

Table 1*Distribution of Gender Classification and Average Rating*


Prompt	Model	Category	Count	Average Rating
p2-1	DeepSeek-V3	F	0	NA
		M	3	4
		U	1	4
	DeepSeek-R1	F	0	NA
		M	0	NA
		U	4	4.25
p2-2	DeepSeek-V3	F	0	NA
		M	6	4.16
		U	0	NA
	DeepSeek-R1	F	0	NA
		M	4	4.25
		U	2	4

First, both models, when they do infer, never produce females. After doing a word-frequency analysis on the ten descriptions generated by DeepSeek-R1 (**Figure 1**), we found the ten most indicative words: *experience, scalable, proven, robust, proficient, passionate, results-driven, user-centric, building, record*. Some of the words — robust, results-driven — appear as the male-encoded words in existing research (Gaucher et al., 2011), while none of these words explicitly indicates gender. However, despite the gender-neutrality in descriptions, neither DeepSeek-R1 nor DeepSeek-V3 classified any description as from a female applicant. This demonstrates that the latent embedding structure aligns “software engineer” summaries with

male representations almost exclusively, regardless of prompt or alignment intensity. Such disproportionality can actually entail bias against females: the model’s behavior does not merely overfit engineer profiles to male, it also implies that “female engineer” lacks representation compatibility with the distributional signals in the résumé summaries.

Figure 1

Most Frequent Words in Descriptions d1 – d10




Second, models with alignment post-training are more risk- and classification-averse under ambiguous instructions. **Figure 2** shows that R1 gives more refusals than V3 when asked to infer from overall semantics (prompt p2-1). R1 refuses all 4 classifications for prompt p2-1, then reduces to 2 out of 6 for p2-2; whereas V3 refuses only 1 out of 4 for p2-1, then none for p2-2. Results of this vertical comparison between DeepSeek-R1 and DeepSeek-V3 shows how stronger reasoning-aligned training makes the model less aggressive in doing classifications, especially when not explicitly

prompted to do so. This can be paralleled to how societal constraints such as promotions for gender equity reduces explicit gender bias. Yet the next findings show that such constraints live only on the surface level.

Figure 2

Vertical Comparison: Effect of Alignment on Bias Recognition




The third finding is key to understanding that alignment post-training affects only how the model handles indeterminacy, not the underlying associative architecture itself. Instructions on making inference “per the description” (p2-1) versus “per the rhetoric, wording, and content” (p2-2) resulted in distinct classification-averse of DeepSeek-R1. **Figure 3** shows that DeepSeek-R1 — though refused all classifications under p2-1 — assigned most of the descriptions to Male under p2-2. The average rating that R1 gives to the Undetermined category is 0.25 lower than that of the Male category under p2-2. The same rise in classification aggressiveness happens to DeepSeek-V3, where the percentage of male classification increased from 75% under p2-1 to 100% under p2-2 (**Table 1**). The prompt-contingency in both R1 and V3 shows that alignment training increases neutrality

of gender perception on only the surface level, and does not mitigate gender bias in LLMs from the underlying representation space.

Figure 3

Prompt-Contingency Observed in DeepSeek-R1



Last, an inspection on the model’s reasoning under prompt p2-2 reveals how gender bias is reproduced through linguistic consolidation. In a case where DeepSeek-R1 does not refuse to classify gender (i.e. classifying the description as from Male), the model justifies its rationale as follows:

“The summary uses neutral professional terminology without any stereotypically gendered markers. The focus on quantifiable results and technical competencies aligns with gender-neutral professional writing conventions. I should emphasize that this inference is based solely on rhetorical patterns.”

The above response demonstrates how the model interprets the description as semantically gender-neutral but rhetorically masculine. This is quite typical of

showing that semantically neutral language can still insist on bias if rhetorically reflecting gender stereotypes.

Building on these findings, we can conclude that while gender bias — inherently tied to stereotypes — is primarily formed through non-verbal symbols, it is ultimately sufficiently reinforced and necessarily sustained through language. We hereby also alert future optimization work of large language models on the limitations of alignment training where it only affects model behavior on a surface level rather than rectifying the representation space.

V. Acknowledgements

This study is constrained by its small sample size and by the fact that both prompts combined hireability ratings with gender inference, which may have unintentionally cued the models to treat gender as evaluation-relevant despite the disclaimer. These limitations do not weaken the central patterns but mark where future work is needed. Future analyses could examine how specific linguistic features — such as assertive versus communal verbs, the frequency of action verbs, hedging, or jargon density — shape gender inference. Additional inquiry is also required on the algorithmic side, where proper redesign or optimization of pre-training corpora and post-training procedures may craft more principled mitigation of gendered representations in large language models.

References

Eckert, P., & McConnell-Ginet, S. (2013). Language and Gender: Introduction to the study of language and gender.

Gaucher, D., Friesen, J., University of Waterloo, Kay, A. C., & Duke University. (2011).

Evidence that gendered wording in job advertisements exists and sustains gender inequality. *Journal of Personality and Social Psychology*, 101–101, 109–128.

<https://doi.org/10.1037/a0022530>

Hadas Koteck, Rikker Dockum, and David Sun. 2023. Gender bias and stereotypes in Large

Language Models. In *Proceedings of The ACM Collective Intelligence Conference (CI '23)*. Association for Computing Machinery, New York, NY, USA, 12–24.

<https://doi.org/10.1145/3582269.3615599>

Ji, M. (2021). Understanding Gender-coded Wording in Job Postings with Word-vectors and

BERT. *Stanford CS224N, Department of East Asian Languages and Cultures, & Stanford University*.

https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1214/reports/final_reports/report009.pdf

Mirza, I., Jafari, A. A., Ozcinar, C., & Anbarjafari, G. (2025). Quantifying Gender Bias in Large

Language Models Using Information-Theoretic and Statistical Analysis. *Information*,

16(5), 358. <https://doi.org/10.3390/info16050358>

Zimman, L. (2017). Transgender language reform: some challenges and strategies for promoting

trans-affirming, gender-inclusive language. *Journal of Language and Discrimination*, 1.1,

84–105. Article. Retrieved from <https://doi.org/10.1558/jld.33139>